

Review Analyzer

NLP Business Case - Automated Customer Reviews Analysis

Group 4

Felipe Doria & Krzysztof Giwojno

IronHack AI Engineering Bootcamp, February 2026

Dashboard

Project 3: NLP Business Case – Group 4, IronHack AI Engineering Bootcamp

TOTAL REVIEWS

28,332

PRODUCTS

65

CATEGORIES

6

AVERAGE RATING

4.51 / 5

Sentiment Distribution



● Positive 89.9% ● Neutral 4.3% ● Negative 5.8%

Rating Distribution



Reviews by Category



Average Rating by Category



We Read Them All!

28,332 reviews. 65 products. One question:
Which Amazon product should I actually buy?

We built a system that:

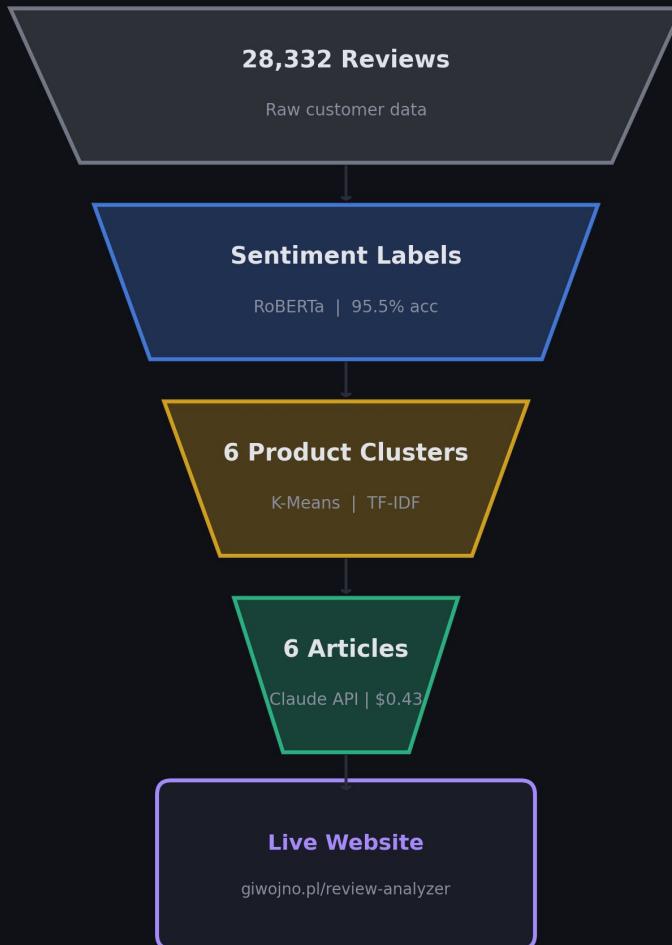
- Reads every review and understands the sentiment
- Groups 65 products into 6 meaningful categories
- Writes recommendation articles - like a tech blogger, but cheaper (\$0.43)

The result? A live website where you can browse categories, compare products, and read AI-generated buying guides.

No backend. No database. Just smart pre-computation.



Pipeline



Introduction

The Real Problem

Companies get thousands of reviews but lack the tools to extract actionable insights at scale.

Our Solution: 3 Models, 1 Pipeline

1. **Classification** feeds sentiment labels to...
2. **Clustering** which groups products, then...
3. **Summarization** writes articles per cluster using the classified reviews as evidence

Each model's output is the next model's input. Classification quality directly affects summarization - if we misclassify a negative review as positive, the AI writes misleading recommendations.

Methodology: Fine-tuned transformers + unsupervised clustering + LLM-powered generation, deployed as a static web app.

Methods

Dataset: Datafiniti Amazon Consumer Reviews (Kaggle)

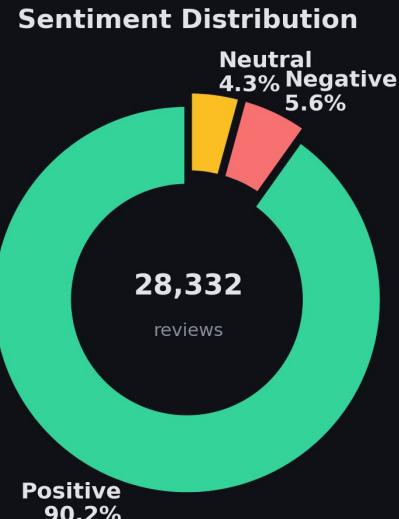
- 28,332 reviews, 65 products, 24 columns
- Reviews from 2009–2019
- All Amazon-branded: Fire Tablets, Kindle, Batteries, Echo

Preprocessing:

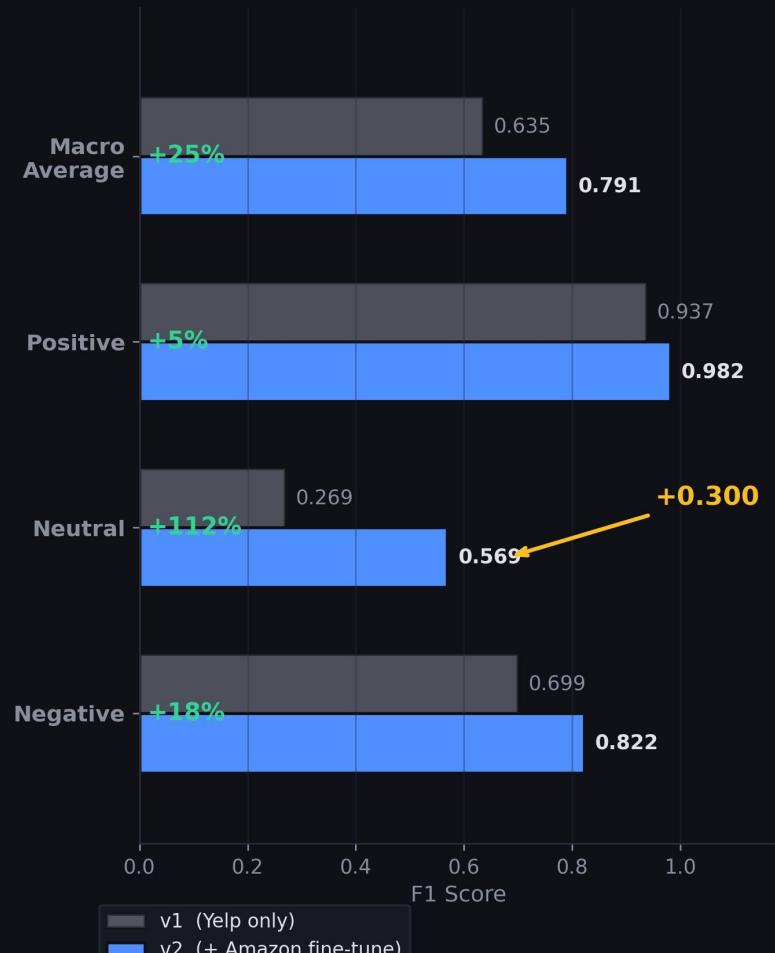
- Dropped 10 columns (>99% null or metadata)
- Kept 35.87% duplicate texts (legitimate short reviews)
- Mapped ratings: 1–2★ → Negative, 3★ → Neutral, 4–5★ → Positive
- No text cleaning for transformers (they handle raw text)

The Challenge: 90% Positive, 6% Negative, 4% Neutral

Training a model on this will predict "Positive" for everything and score 90% accuracy - but that's useless.



Classification F1: v1 vs v2



Model 1 — Classification

Task: Classify each review as Positive, Negative, or Neutral

What we tried:

- **RoBERTa-base** (125M params) — transformer pre trained on large text corpora

Two-stage training:

1. Pre-trained on 650K Yelp reviews (3-class) → learns general sentiment
2. Fine-tuned on Amazon data with **class weights** → adapts to our domain

Key technique: Class Weights

- Neutral gets 7.8x more importance during training
- Negative gets 6.0x
- Positive gets 0.4x (**downweighted**)

This forces the model to pay attention to rare classes instead of just predicting "Positive" every time.

Confusion Matrix (v2 Test Set)

Model 1 - Evaluation

95.48% Accuracy | 0.791 Macro F1

Positive: Precision 0.985 | Recall 0.979 | F1 0.982

Neutral: Precision 0.550 | Recall 0.589 | F1 0.569

Negative: Precision 0.804 | Recall 0.842 | F1 0.822

Only 256 errors out of 5,667 test reviews (4.5%)

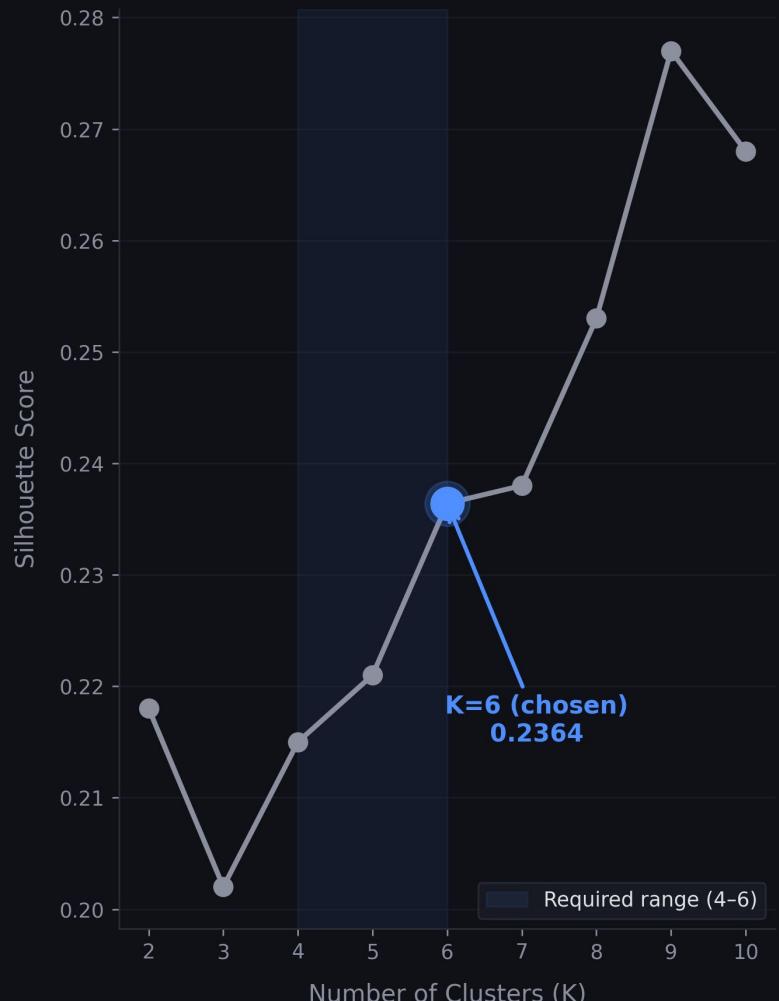
Where it fails:

- **Positive** ↔ **Neutral** boundary (most errors)
- **Sarcasm and mixed-sentiment reviews**
- "It's okay I guess" — is that Neutral or Positive?

The worst error (**Negative** → **Positive**) happens only 12 times



Silhouette Score by K



Model 2 — Clustering

Task: Group 65 products into 4–6 meaningful categories

Why not use existing categories?

- **primaryCategories** is useless — 92% of products are just "Electronics" or "Health & Beauty"
- We needed categories a shopper would recognize

Method:

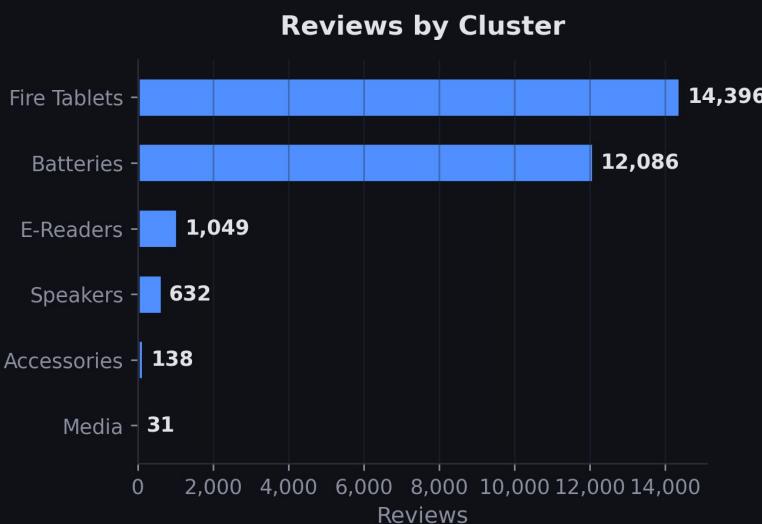
- Combined product names + cleaned category labels
- TF-IDF vectorization (unigrams + bigrams, 200 features)
- K-Means clustering
- Tested K=2 to K=10

Chose K=6 - best silhouette within required 4-6 range (0.2364)

Model 2 - Evaluation

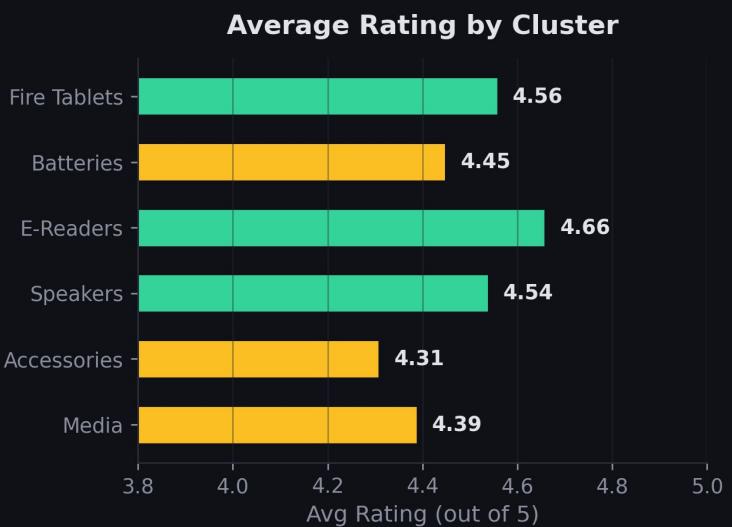
6 Clusters Found

Fire Tablets		20 products		14,396 reviews		★ 4.56
Batteries		7 products		12,086 reviews		★ 4.45
E-Readers		10 products		1,049 reviews		★ 4.66
Smart Speakers		9 products		632 reviews		★ 4.54
Accessories		11 products		138 reviews		★ 4.31
Media & Home		8 products		31 reviews		★ 4.39



Do they make sense? Yes:

- Every Kindle → E-Readers ✓
- Every Fire tablet variant → Fire Tablets ✓
- Batteries separated from electronics ✓
- Named clusters manually after inspecting members



Product Clusters (TF-IDF + PCA)



API Output (Claude Sonnet)

The Fire HD 8 with Alexa stands out as the category leader with 2,445 reviews and a 4.6 rating. Users consistently praise the value for money, screen quality, and Alexa integration.

However, app limitations and occasional slowness are recurring complaints. For budget-conscious buyers seeking a reliable entertainment tablet, the Fire HD 8 remains the top recommendation.

\$0.43

vs

Local Output (Flan-T5-base)

Fire HD 8 Tablet with Alexa has rating 4.6 from 2445 reviews.

\$0.00

Same data. Same task.

250M params vs 70B+ params

Model 3 - Summarization

Task: Generate recommendation articles per category

Two parallel approaches:

API - Claude Sonnet

- Sample 5 positive, 5 negative, 3 neutral reviews per product
- Send with stats to LLM via structured prompts
- Output: 600-word articles with top products, complaints, recommendations
- **Cost:** \$0.43 total for 39 summaries

Local - Flan-T5-base (250M params)

- Build evidence briefs with product rankings + TF-IDF complaints
- Smart ranking: `avg_rating × log(review_count)`
- Run entirely on GPU, no API needed
- **Cost:** \$0.00

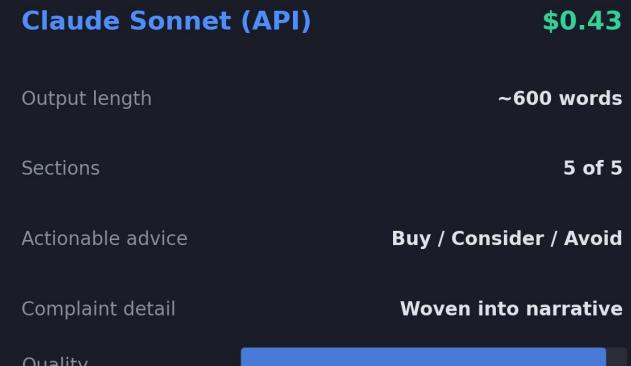
Model 3 - Evaluation

How Do You Evaluate Generated Text?

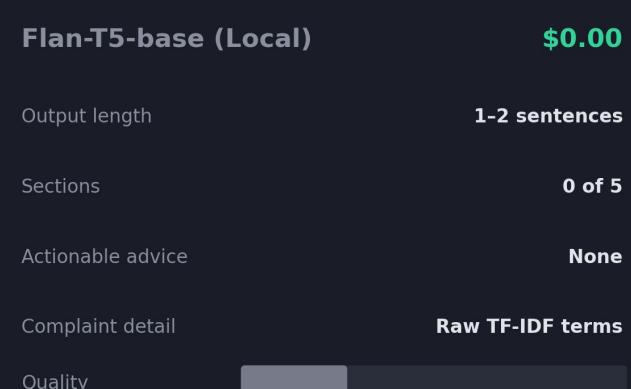
	API (Claude)	Local (T5)
Structure	All sections	No sections
Length	~600 words	1–2 sentences
Factual accuracy	Correct stats	Correct but minimal
Actionable advice	Buy/Consider	None
Complaints	In narrative	Raw TF-IDF terms
Cost	\$0.43	\$0.00

Both approaches deployed to the website side-by-side
Users can switch between API and Local tabs to compare
Demonstrates the quality-cost tradeoff transparently

<https://giwojno.pl/review-analyzer/>



vs



Model 3 - Evaluation

How Do You Evaluate Generated Text?

	API (Claude)	Local (T5)
Structure	All sections	No sections
Length	~600 words	1–2 sentences
Factual accuracy	Correct stats	Correct but minimal
Actionable advice	Buy/Consider	None
Complaints	In narrative	Raw TF-IDF terms
Cost	\$0.43	\$0.00

Both approaches deployed to the website side-by-side
Users can switch between API and Local tabs to compare
Demonstrates the quality-cost tradeoff transparently

<https://giwojno.pl/review-analyzer/>



Categories — API Model Output

Generated by Claude Sonnet (claude-sonnet-4-20250514) via the Anthropic API. Each category summary is a full recommendation article (~600 word top products, common complaints, worst product, and buying advice. Total cost: \$0.43 USD for 39 API calls. Temperature: 0.3 for consistent, factual out-

[← Back to categories](#)

Fire Tablets

20 products 14,396 reviews ★ 4.56

Generated by Claude Sonnet (claude-sonnet-4-20250514) via Anthropic API

Fire Tablets: A Comprehensive Buying Guide Based on Customer Reviews

Category Overview

Fire Tablets represent Amazon's budget-friendly entry into the tablet market, with 20 different models earning an impressive 4.56/5 average rating across 14,396 customer reviews. The category shows strong customer satisfaction, with most models achieving over 90% positive reviews, making them a compelling choice for budget-conscious consumers seeking basic tablet functionality.

Top 3 Recommended Products

1. Fire HD 8 Tablet with Alexa (16GB, Tangerine)

With 2,443 reviews and a 4.6/5 rating, this model leads in both popularity and satisfaction. Customers praise its excellent value at around \$89, calling it "fast at everything except gameplay" with "great stereo speakers." The 8" HD display receives consistent praise for movies and reading, though some note it struggles with direct sunlight.

2. All-New Fire HD 8 Tablet (16GB)

Nearly matching the top model with 2,370 reviews and 4.58/5 rating, customers highlight its superior screen quality compared to the 7" models and adequate performance for web browsing and streaming. Parents particularly appreciate the FreeTime parental controls available across Fire models.

3. Fire Kids Edition Tablet (7" Display, Pink)

Leading the kids' category with 1,676 reviews and 4.53/5 rating, this tablet stands out for its exceptional durability. Parents rave about the "thick, heavy duty" protective case that withstands drops, bites, and rough handling. The 2-year no-questions-asked replacement warranty, including cracked screens, receives universal praise.

Categories — Local Model Output

Generated by Flan-T5-base (google/flan-t5-base) — a 250M parameter seq2seq transformer running locally on GPU (NVIDIA RTX 4050). Unlike the API-based summaries, this model runs entirely offline with no API costs. The smaller model size results in shorter, more factual outputs compared to the Claude Sonnet API approach.

[← Back to categories](#)

Fire Tablets

20 products 14,396 reviews ★ 4.56

Generated by Flan-T5-base (250M params) running locally on GPU

Fire Tablets

The Fire HD 8 Tablet with Alexa, 8 HD Display, 16 GB, Tangerine – with Special Offers has rating 4.6 from 2445 reviews. The All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi, 16 GB – Includes Special Offers, Black has rating 4.58 from 2370 reviews. The lowest rated is Fire Tablet, 7 Display, Wi-Fi, 16 GB – Includes Special Offers, Black with rating 4.51 from 1028 reviews.

Top Picks

Fire HD 8 Tablet with Alexa, 8 HD Display, 16 GB, Tangerine – with Special Offers

Rating: 4.6 • Reviews: 2445

Complaints: tablet; just; good; bought; amazon

All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi, 16 GB – Includes Special Offers, Black

Rating: 4.58 • Reviews: 2370

Complaints: tablet; just; good; amazon; apps

Fire Kids Edition Tablet, 7 Display, Wi-Fi, 16 GB, Pink Kid-Proof Case

Rating: 4.53 • Reviews: 1702

Complaints: tablet; charge; apps; slow; good

Takeaway

Try It Live

Scan to open on your phone



What we built: A complete NLP pipeline from raw reviews → deployed web app

- **95.48% accuracy → 6 clusters → 39 summaries → Live website**
- **Both API (\$0.43) + Local models deployed side-by-side**
- Live at <https://giwojno.pl/review-analyzer>

What We Learned

- **Class weights matter** - they doubled Neutral F1 from 0.27 to 0.57
- **Model size matters** - 250M params can't do what an API model does for \$0.43
- **Keep it simple** - static site, no backend, pre-compute everything

Bonus point pitch: We deployed a [fully functional public website](#) on real hosting - not localhost, not a notebook, not Gradio. A static site architecture that pre-computes all results, needs zero backend, and works on the cheapest shared hosting available. You can open it [right now](#) on your phone.

giwojno.pl/review-analyzer

Takeaway

What we built: A complete NLP pipeline from raw reviews → deployed web app

- **95.48% classification accuracy (RoBERTa)**
- **6 meaningful product clusters (K-Means)**
- **39 AI-generated summaries for \$0.43**
- **Both API + Local models deployed side-by-side**
- Live at <https://giwojno.pl/review-analyzer>

Bonus point pitch: We deployed a **fully functional public website** on real hosting - not localhost, not a notebook, not Gradio. A static site architecture that pre-computes all results, needs zero backend, and works on the cheapest shared hosting available. You can open it **right now** on your phone.

Try It Live

Scan to open on your phone



giwojno.pl/review-analyzer

Demo

Dashboard

Project 3: NLP Business Case — Group 4, IronHack AI Engineering Bootcamp

TOTAL REVIEWS

28.332

PRODUCTS

65

CATEGORIES

6

AVERAGE RATING

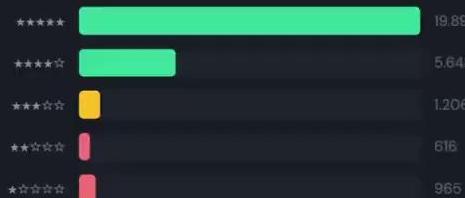
4.51 / 5

Sentiment Distribution



● Positive 89.9% ● Neutral 4.3% ● Negative 5.8%

Rating Distribution



Reviews by Category

Fire Tablets 14.398

Batteries & Household 12.086

E-Readers 1.049

Smart Speakers 632

Average Rating by Category

Fire Tablets 4.56/5

Batteries & Household 4.45/5

E-Readers 4.66/5

Smart Speakers 4.54/5



QUESTIONS?

